

input_ids : 词元化的句子，每个词元用id表示

B是batch_size，L是句子长度

input_ids作为输入送入embedding层

> `self.embedding = nn.Embedding(vocab_size, d_model, **factory_kwargs)`

vocab_size是词汇表的大小，d_model表示状态空间维数（词向量维数）


接下来送入若干个Block（ResidualBlock, named from mamba-minimal）

得到的结果作归一化（RMSNorm or nn.LayerNorm）后送入线性层得到logits



logits可以用于classify或者decode之后用于generate

```
1 class MixerModel(nn.Module):
2     def __init__(...) -> None:
3         factory_kwargs = {"device": device, "dtype": dtype}
4         super().__init__()
5         self.residual_in_fp32 = residual_in_fp32
6
7         self.embedding = nn.Embedding(vocab_size, d_model, **factory_kwargs)
8
9         self.fused_add_norm = fused_add_norm
10        if self.fused_add_norm:
11            if layer_norm_fn is None or rms_norm_fn is None:
12                raise ImportError("Failed to import Triton LayerNorm / RMSNorm kernels")
13
14        self.layers = nn.ModuleList( [ create_block(...) for i in range(n_layer) ] )
15
16        self.norm_f = (nn.LayerNorm if not rms_norm else RMSNorm)(
17            d_model, eps=norm_epsilon, **factory_kwargs
18        )
19
20        self.apply(
21            partial(
22                _init_weights,
23                n_layer=n_layer,
24                **(initializer_cfg if initializer_cfg is not None else {}),
25            )
26        )
```

(B,L)



(B,L,d_model)



(B,L,d_model)



(B,L,vocab_size)

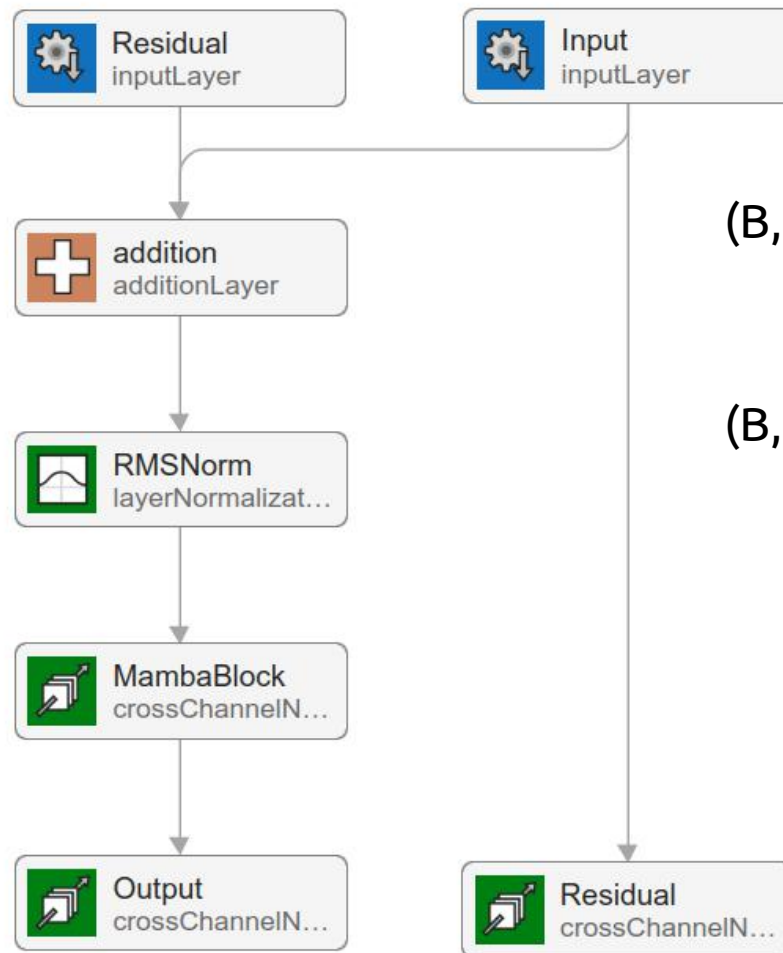


d_{inner} 是内部维度， $d_{\text{inner}}=d_{\text{model}} * \text{expand}$ ($\text{expand}=2$ by default)
 在MambaBlock内部，输入经过第一个线性层之后等分为 d_{inner} 的两份，
 分别交付两条计算路线。F.silu是激活函数。

```

1 class Mamba(nn.Module):
2     def __init__(...):
3         factory_kwargs = {"device": device, "dtype": dtype}
4         super().__init__()
5         self.d_model = d_model
6         self.d_state = d_state
7         self.d_conv = d_conv
8         self.expand = expand
9         self.d_inner = int(self.expand * self.d_model)
10        ...
11        self.use_fast_path = use_fast_path
12        self.layer_idx = layer_idx
13
14        self.in_proj = nn.Linear(
15            self.d_model,
16            self.d_inner * 2,
17            bias=bias,
18            **factory_kwargs
19        )
20
21        self.conv1d = nn.Conv1d(
22            in_channels=self.d_inner,
23            out_channels=self.d_inner,
24            bias=conv_bias,
25            kernel_size=d_conv,
26            groups=self.d_inner,
27            padding=d_conv - 1,
28            **factory_kwargs,
29        )
30
31        self.activation = "silu"
32        self.act = nn.SiLU()
33
34        ...
35
36        self.out_proj = nn.Linear(
37            self.d_inner,
38            self.d_model,
39            bias=bias,
40            **factory_kwargs
41        )
    
```

inside ResidualBlock



inside MambaBlock

